

Logistic Regression for Predictive Modeling

Data preparation

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sn
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.model_selection import cross_val_predict, train_test_split
from sklearn import metrics
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: df = pd.read_csv('/Users/ebeth/Desktop/Churn Data/churn_clean.csv')
```

```
In [3]: df.head()
```

```
Out[3]:
```

	CaseOrder	Customer_id	Interaction	UID	City	Sta
0	1	K409198	aa90260b-4141-4a24-8e36-b04ce1f4f77b	e885b299883d4f9fb18e39c75155d990	Point Baker	,
1	2	S120509	fb76459f-c047-4a9d-8af9-e0f7d4ac2524	f2de8bef964785f41a2959829830fb8a	West Branch	
2	3	K191035	344d114c-3736-4be5-98f7-c72c281e2d35	f1784cfa9f6d92ae816197eb175d3c71	Yamhill	(
3	4	D90850	abfa2b40-2d43-4994-b15a-989b8c79e311	dc8a365077241bb5cd5ccd305136b05e	Del Mar	(
4	5	K662701	68a861fd-0d20-4e51-a587-8a90407ee574	aabb64a116e83fdc4befc1fbab1663f9	Needville	.

5 rows x 50 columns

```
In [4]: df.shape
```

```
Out[4]: (10000, 50)
```

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 50 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   CaseOrder                            10000 non-null  int64
1   Customer_id                          10000 non-null  object
2   Interaction                          10000 non-null  object
3   UID                                  10000 non-null  object
4   City                                 10000 non-null  object
5   State                                10000 non-null  object
6   County                              10000 non-null  object
7   Zip                                  10000 non-null  int64
8   Lat                                  10000 non-null  float64
9   Lng                                  10000 non-null  float64
10  Population                           10000 non-null  int64
11  Area                                 10000 non-null  object
12  TimeZone                            10000 non-null  object
13  Job                                  10000 non-null  object
14  Children                            10000 non-null  int64
15  Age                                  10000 non-null  int64
16  Income                              10000 non-null  float64
17  Marital                             10000 non-null  object
18  Gender                              10000 non-null  object
19  Churn                               10000 non-null  object
20  Outage_sec_perweek                  10000 non-null  float64
21  Email                               10000 non-null  int64
22  Contacts                            10000 non-null  int64
23  Yearly_equip_failure                10000 non-null  int64
24  Techie                              10000 non-null  object
25  Contract                            10000 non-null  object
26  Port_modem                          10000 non-null  object
27  Tablet                              10000 non-null  object
28  InternetService                    10000 non-null  object
29  Phone                               10000 non-null  object
30  Multiple                            10000 non-null  object
31  OnlineSecurity                      10000 non-null  object
32  OnlineBackup                        10000 non-null  object
33  DeviceProtection                    10000 non-null  object
34  TechSupport                         10000 non-null  object
35  StreamingTV                         10000 non-null  object
36  StreamingMovies                     10000 non-null  object
37  PaperlessBilling                    10000 non-null  object
38  PaymentMethod                       10000 non-null  object
39  Tenure                              10000 non-null  float64
40  MonthlyCharge                       10000 non-null  float64
41  Bandwidth_GB_Year                  10000 non-null  float64
42  Item1                              10000 non-null  int64
43  Item2                              10000 non-null  int64
44  Item3                              10000 non-null  int64
45  Item4                              10000 non-null  int64
46  Item5                              10000 non-null  int64
47  Item6                              10000 non-null  int64
48  Item7                              10000 non-null  int64
49  Item8                              10000 non-null  int64
dtypes: float64(7), int64(16), object(27)
memory usage: 3.8+ MB
```

```
In [6]: df2 = df.drop(['CaseOrder', 'Customer_id', 'Interaction', 'UID', 'Lat', 'Lng', 'Ti
```

Take a closer look at state and zip code to see if they will be included in the regression. Looking at the number of unique values will reveal if using these variables will make the regression equation unwieldy.

```
In [7]: df2.nunique(axis = 0)
```

```
Out[7]: State                52
Zip                8583
Population         5933
Area                3
Children           11
Age                72
Income            9993
Marital            5
Gender             3
Churn              2
Outage_sec_perweek 9986
Email             23
Contacts           8
Yearly_equip_failure 6
Techie            2
Contract           3
Port_modem         2
Tablet             2
InternetService    3
Phone              2
Multiple           2
OnlineSecurity     2
OnlineBackup       2
DeviceProtection   2
TechSupport        2
StreamingTV        2
StreamingMovies    2
PaperlessBilling   2
PaymentMethod      4
Tenure             9996
MonthlyCharge      750
Bandwidth_GB_Year 10000
Item1              7
Item2              7
Item3              8
Item4              7
Item5              7
Item6              8
Item7              7
Item8              8
dtype: int64
```

There are 52 states and 8583 unique zip codes in the data set. For this overall look at churn this many variables would not be helpful. Exploring individual states or zip codes at a later date might be beneficial. State and zip will be dropped from the data set.

```
In [8]: df3 = df2.drop(['State', 'Zip'], axis = 1)
```

Look for missing values, duplicates, and outliers

```
In [9]: df3.isnull().count()
```

```
Out[9]: Population      10000
Area      10000
Children  10000
Age      10000
Income    10000
Marital   10000
Gender    10000
Churn      10000
Outage_sec_perweek  10000
Email      10000
Contacts  10000
Yearly_equip_failure  10000
Techie     10000
Contract   10000
Port_modem 10000
Tablet     10000
InternetService  10000
Phone      10000
Multiple   10000
OnlineSecurity  10000
OnlineBackup  10000
DeviceProtection  10000
TechSupport  10000
StreamingTV  10000
StreamingMovies  10000
PaperlessBilling  10000
PaymentMethod  10000
Tenure     10000
MonthlyCharge  10000
Bandwidth_GB_Year  10000
Item1      10000
Item2      10000
Item3      10000
Item4      10000
Item5      10000
Item6      10000
Item7      10000
Item8      10000
dtype: int64
```

```
In [10]: df3.duplicated()
```

```
Out[10]: 0      False
1      False
2      False
3      False
4      False
...
9995   False
9996   False
9997   False
9998   False
9999   False
Length: 10000, dtype: bool
```

```
In [11]: df3.describe()
```

```
Out[11]:
```


	Population	Children	Age	Income	Outage_sec_perweek	
count	10000.000000	10000.0000	10000.000000	10000.000000	10000.000000	10000.
mean	9756.562400	2.0877	53.078400	39806.926771	10.001848	12
std	14432.698671	2.1472	20.698882	28199.916702	2.976019	3.
min	0.000000	0.0000	18.000000	348.670000	0.099747	1.
25%	738.000000	0.0000	35.000000	19224.717500	8.018214	10.
50%	2910.500000	1.0000	53.000000	33170.605000	10.018560	12.
75%	13168.000000	3.0000	71.000000	53246.170000	11.969485	14.
max	111850.000000	10.0000	89.000000	258900.700000	21.207230	23.

Univariate Visualizations

Target variable Churn

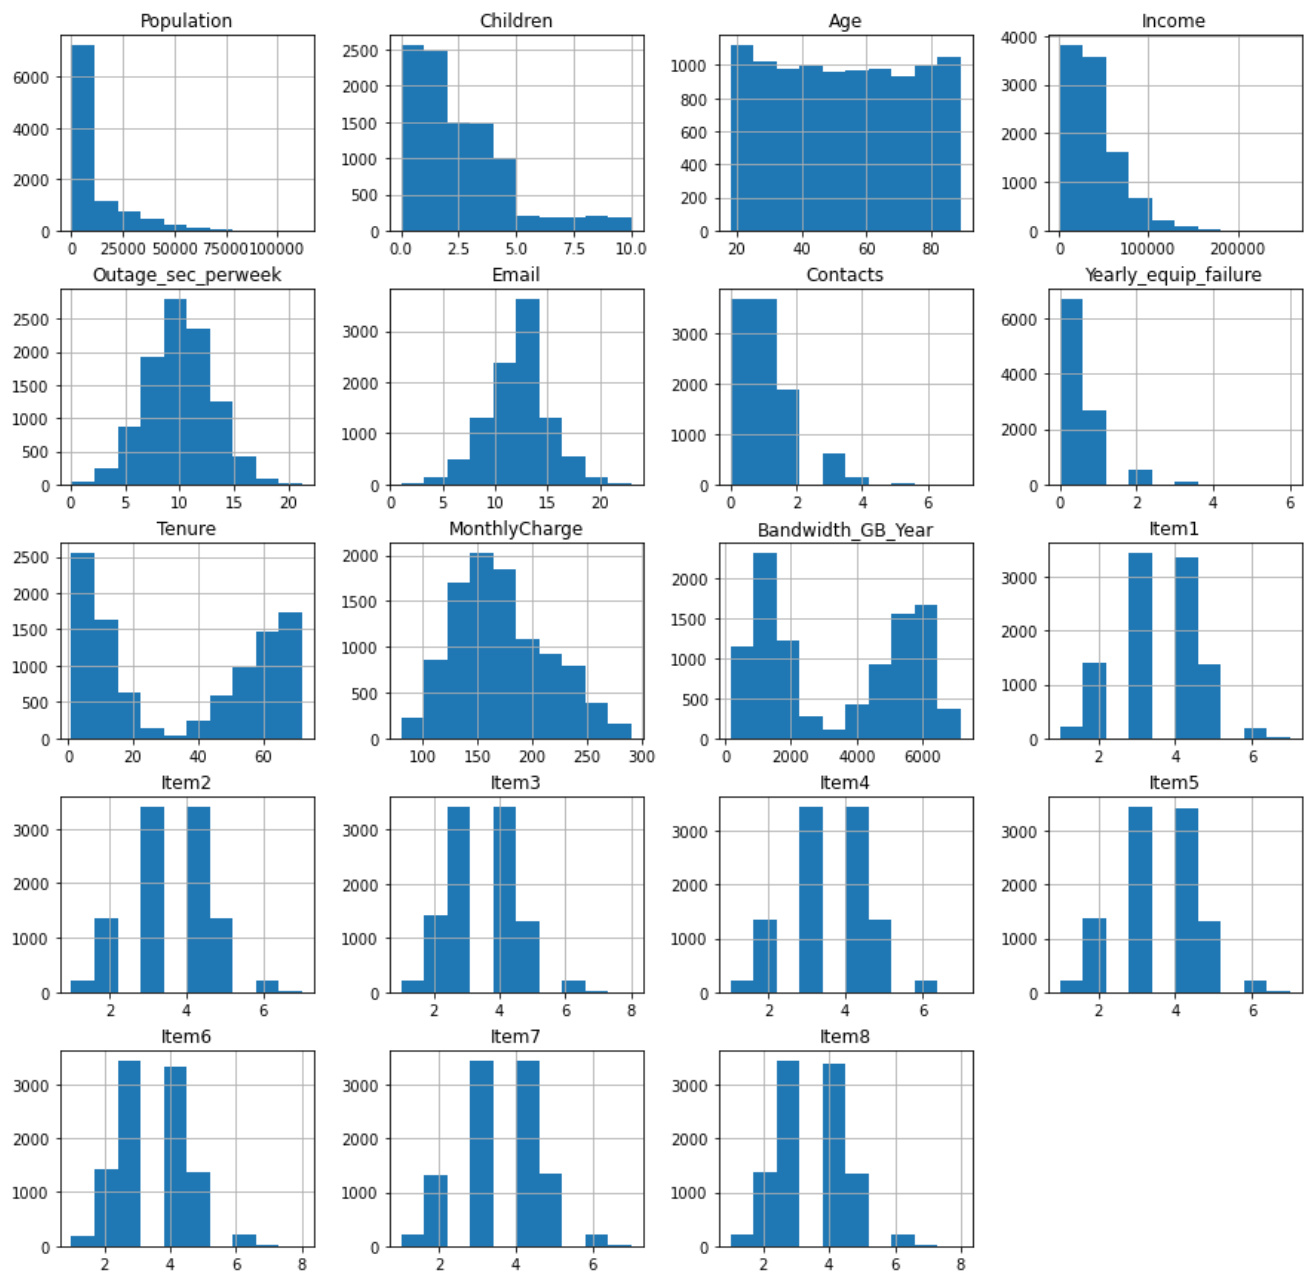
```
In [12]: df3['Churn'].value_counts().plot.bar(title = 'Churn')
```

```
Out[12]: <AxesSubplot:title={'center':'Churn'}>
```




Predictor Variables


```
In [13]: df3.hist(figsize = (15,15));
```




```
In [14]: df3['Area'].value_counts().plot.bar(title = 'Area Types');
```




```
In [15]: df3['Marital'].value_counts().plot.bar(title = 'Marital Status');
```




```
In [16]: df3['Gender'].value_counts().plot.bar(title = 'Gender');
```




```
In [17]: df3['Contract'].value_counts().plot.bar(title = 'Length of Contract');
```



```
In [18]: df3['PaymentMethod'].value_counts().plot.bar(title = 'Method of Payment');
```


```
In [19]: df3['InternetService'].value_counts().plot.bar(title = 'Internet Types');
```



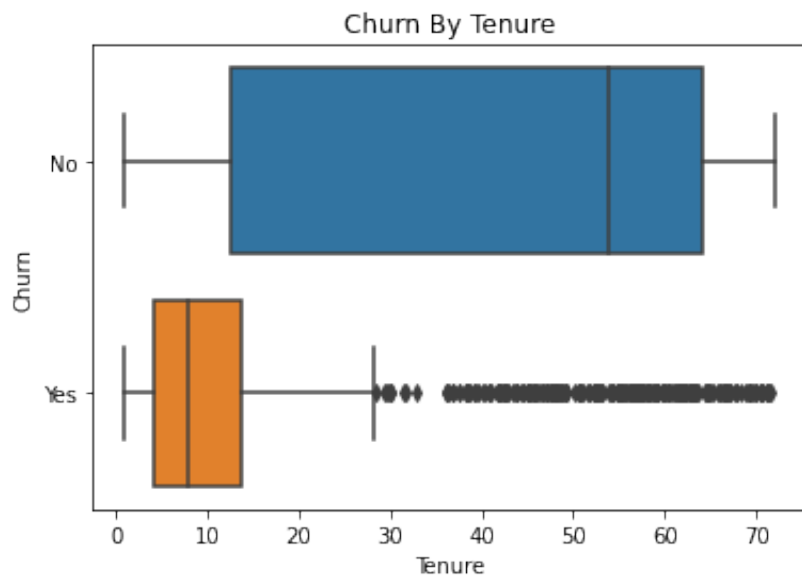
Bivariate Visualizations

<https://seaborn.pydata.org/tutorial/categorical.html>


```
In [20]: sn.boxplot(df3['Age'], df3['Churn']).set_title('Churn By Age');
```




```
In [21]: sn.boxplot(df3['Tenure'], df3['Churn']).set_title('Churn By Tenure');
```




```
In [22]: sn.boxplot(df3['MonthlyCharge'], df3['Churn']).set_title('Churn By Monthly Ch');
```




```
In [23]: sn.boxplot(df3['Income'], df3['Churn']).set_title('Churn By Income');
```



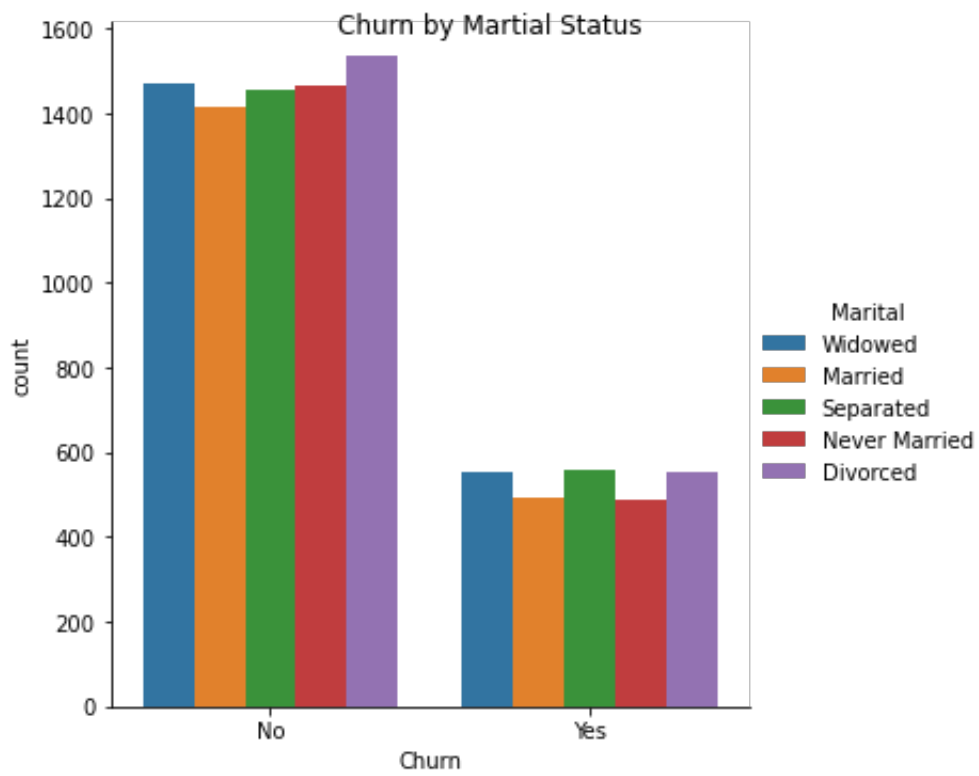
```
In [24]: sn.boxplot(df3['Outage_sec_perweek'], df3['Churn']).set_title('Churn By Outag
```




```
In [25]: sn.boxplot(df3['Bandwidth_GB_Year'], df3['Churn']).set_title('Churn By Bandwi
```



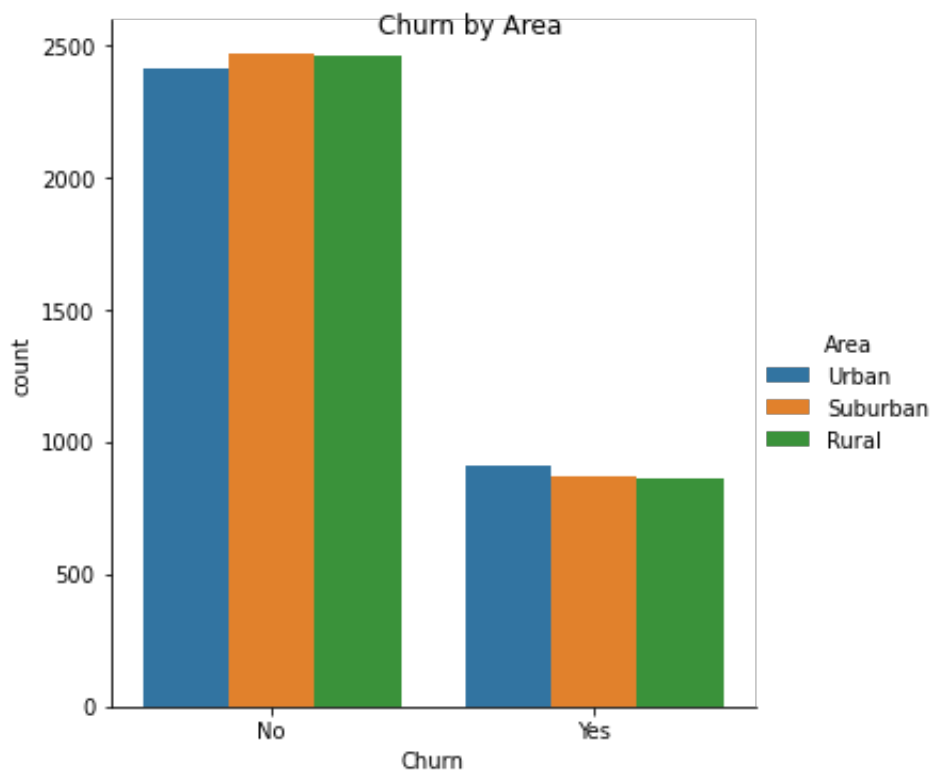
```
In [26]: g = sn.catplot(x='Churn', hue="Marital", kind="count", data=df3);
g.fig.suptitle('Churn by Martial Status');
```



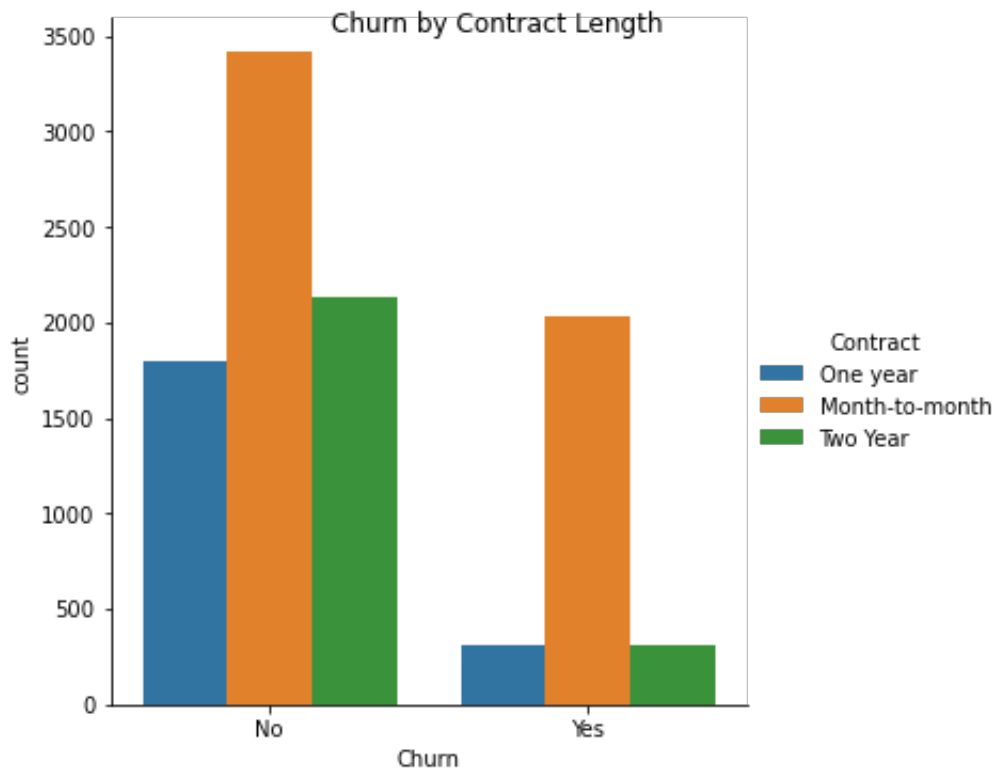
```
In [27]: g = sns.catplot(x='Churn', hue="Gender", kind="count", data=df3);
g.fig.suptitle('Churn by Gender');
```



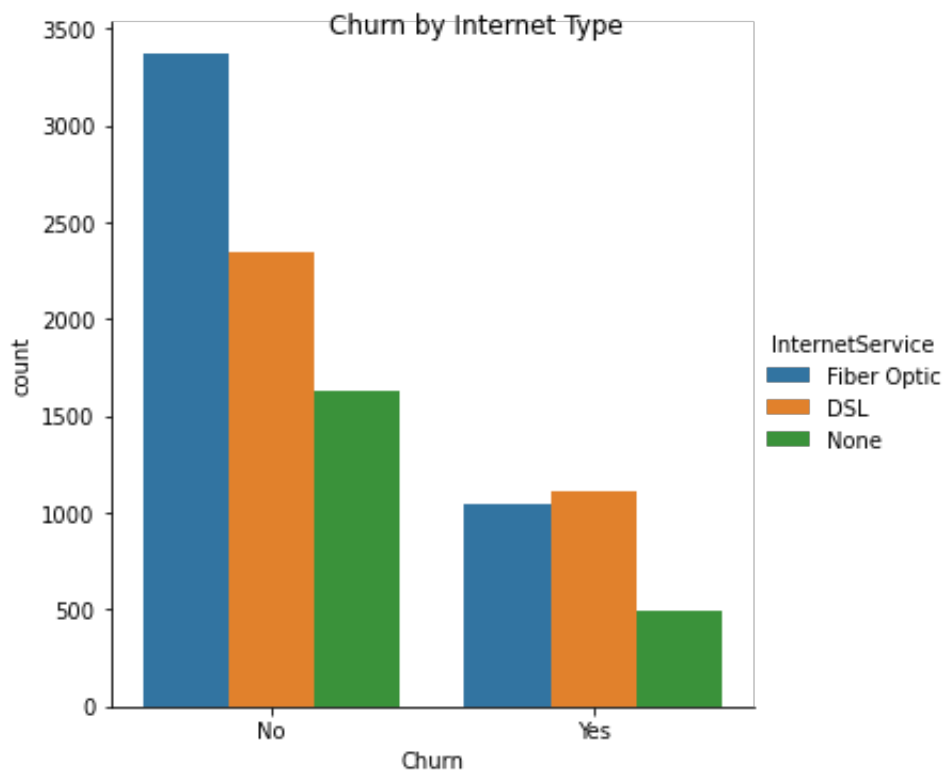
```
In [28]: g = sns.catplot(x='Churn', hue="Area", kind="count", data=df3);
g.fig.suptitle('Churn by Area');
```



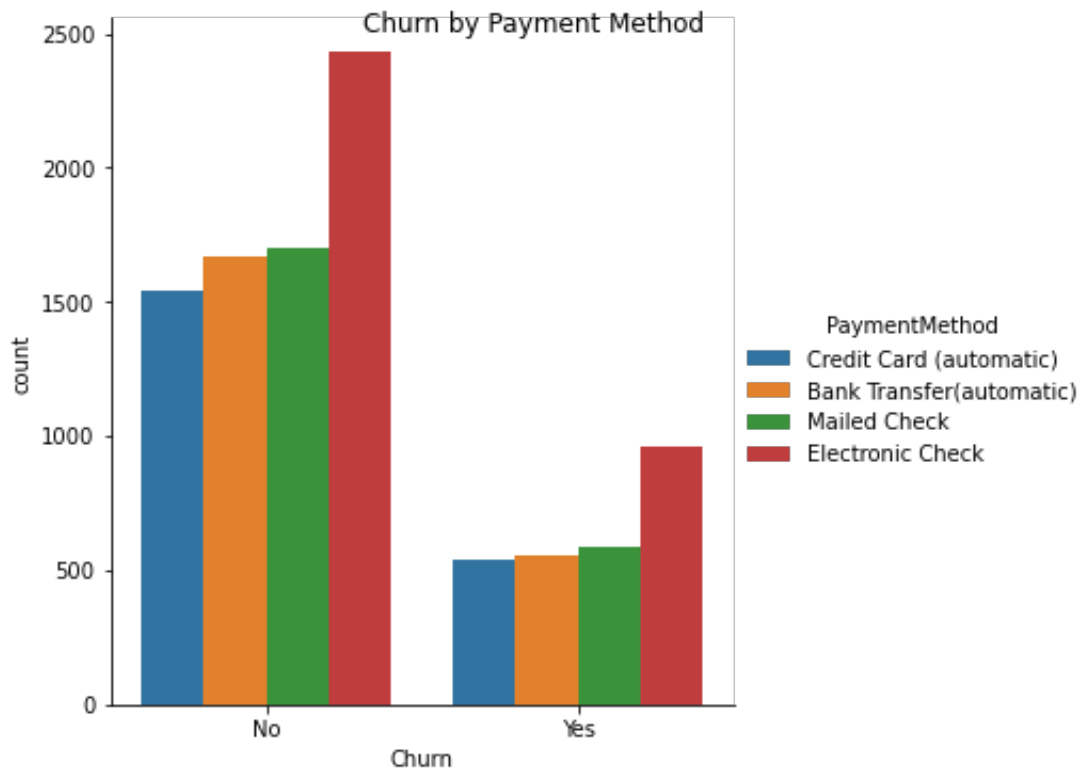
```
In [29]: g = sns.catplot(x='Churn', hue="Contract", kind="count", data=df3);  
g.fig.suptitle('Churn by Contract Length');
```



```
In [30]: g = sns.catplot(x='Churn', hue="InternetService", kind="count", data=df3);  
g.fig.suptitle('Churn by Internet Type');
```



```
In [31]: g = sns.catplot(x='Churn', hue="PaymentMethod", kind="count", data=df3);
g.fig.suptitle('Churn by Payment Method');
```



Create Dummy variables for all categorical columns and drop unneeded columns. (code used from: <https://towardsdatascience.com/the-dummys-guide-to-creating-dummy-variables-f21faddb1d40>)

```
In [32]: df3 = pd.get_dummies(df3)
df3 = df3.drop(columns = ['Churn_No', 'Techie_No', 'Port_modem_No', 'Tablet_No'])
```

```
In [33]: df3.head()
```

```
Out[33]:
```

	Population	Children	Age	Income	Outage_sec_perweek	Email	Contacts	Yearly equip_f
0	38	0	68	28561.99	7.978323	10	0	
1	10446	1	27	21704.77	11.699080	12	0	
2	3735	4	50	9609.57	10.752800	9	0	
3	13863	1	48	18925.23	14.913540	15	2	
4	11352	0	83	40074.19	8.147417	16	2	

5 rows × 53 columns

Save Copy of Clean Data

```
In [34]: df3.to_csv('logistic_prepared_churn.csv')
```

Construct Initial Logistic Regression Model

<https://towardsdatascience.com/logistic-regression-using-python-sklearn-numpy-mnist-handwriting-recognition-matplotlib-a6b31e2b166a>

```
In [35]: y = df3['Churn_Yes']
X = df3.loc[:, df3.columns != 'Churn_Yes']
```

```
In [36]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.80, r
```

```
In [37]: allmodel = LogisticRegression(solver = 'liblinear', random_state = 0)
```

```
In [38]: allmodel.fit(X_train, y_train)
```


```
Out[38]: LogisticRegression(random_state=0, solver='liblinear')
```

```
In [39]: allpredictions = allmodel.predict(X_test)
print(allpredictions)
```

```
[0 0 1 ... 0 1 0]
```

<https://www.datacamp.com/community/tutorials/understanding-logistic-regression-python>



```
In [40]: cm = metrics.confusion_matrix(y_test, allpredictions)
plt.figure(figsize = (15,15))
sn.heatmap(cm, annot = True, fmt = 'g');
plt.ylabel('Actual');
plt.xlabel('Predicted');
```



```
In [41]: accuracy = print('Accuracy:', metrics.accuracy_score(y_test, allpredictions))
precision = print('Precision:', metrics.precision_score(y_test, allpredictions))
recall = print('Recall:', metrics.recall_score(y_test, allpredictions))
```


Accuracy: 0.859125
Precision: 0.7690217391304348
Recall: 0.6683986773736419

```
In [42]: pred_prob = allmodel.predict(X_test)
fpr, tpr, _ = metrics.roc_curve(y_test, pred_prob)
auc = metrics.roc_auc_score(y_test, pred_prob)
plt.plot(fpr, tpr, label = 'auc = '+str(auc))
plt.legend(loc = 4)
plt.show()
```



Model Reduction

```
In [43]: plt.figure(figsize = (30,30))
heat_map = sn.heatmap(df3.corr(),vmin = -1, vmax = 1, annot = True, fmt = '.2')
heat_map.set_title('Correlation Heatmap');
```



<https://datascience.stackexchange.com/questions/39137/how-can-i-check-the-correlation-between-features-and-target-variable>

In [44]:

```
pd.options.display.max_rows = 100
df3[df3.columns[1:]].corr()['Churn_Yes'][:]
```

```

Out[44]: Children -0.004264
Age 0.005630
Income 0.005937
Outage_sec_perweek -0.000156
Email 0.012326
Contacts 0.008567
Yearly_equip_failure -0.015927
Tenure -0.485475
MonthlyCharge 0.372938
Bandwidth_GB_Year -0.441669
Item1 -0.007341
Item2 -0.013253
Item3 -0.011143
Item4 -0.003396
Item5 -0.013971
Item6 0.001130
Item7 -0.008851
Item8 0.005653
Area_Rural -0.008971
Area_Suburban -0.006574
Area_Urban 0.015554
Marital_Divorced -0.000769
Marital_Married -0.007731
Marital_Never Married -0.017331
Marital_Separated 0.014854
Marital_Widowed 0.010622
Gender_Female -0.027021
Gender_Male 0.028061
Gender_Nonbinary -0.003341
Churn_Yes 1.000000
Techie_Yes 0.066722
Contract_Month-to-month 0.267653
Contract_One year -0.139043
Contract_Two Year -0.178337
Port_modem_Yes 0.008157
Tablet_Yes -0.002779
InternetService_DSL 0.093487
InternetService_Fiber Optic -0.058472
InternetService_None -0.037742
Phone_Yes -0.026297
Multiple_Yes 0.131771
OnlineSecurity_Yes -0.013540
OnlineBackup_Yes 0.050508
DeviceProtection_Yes 0.056489
TechSupport_Yes 0.018838
StreamingTV_Yes 0.230151
StreamingMovies_Yes 0.289262
PaperlessBilling_Yes 0.007030
PaymentMethod_Bank Transfer(automatic) -0.017795
PaymentMethod_Credit Card (automatic) -0.006693
PaymentMethod_Electronic Check 0.029914
PaymentMethod_Mailed Check -0.009626
Name: Churn_Yes, dtype: float64

```

Reduced Model

```

In [45]: df4 = df3[['Churn_Yes', 'Tenure', 'MonthlyCharge', 'Bandwidth_GB_Year', 'Cont

```

```
In [46]: y = df4['Churn_Yes']  
X = df4.loc[:, df4.columns != 'Churn_Yes']
```

```
In [47]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.80, r
```

```
In [48]: redmodel = LogisticRegression(solver = 'liblinear', random_state = 0)
```


```
In [49]: redmodel.fit(X_train, y_train)
```

```
Out[49]: LogisticRegression(random_state=0, solver='liblinear')
```

```
In [50]: redpredictions = redmodel.predict(X_test)  
print(redpredictions)
```

```
[0 0 1 ... 0 1 0]
```


```
In [51]: cm = metrics.confusion_matrix(y_test, redpredictions)  
plt.figure(figsize = (15,15))  
sn.heatmap(cm, annot = True, fmt = 'g');  
plt.ylabel('Actual');  
plt.xlabel('Predicted');
```



```
In [52]: accuracy = print('Accuracy:', metrics.accuracy_score(y_test, redpredictions))
precision = print('Precision:', metrics.precision_score(y_test, redpredictions)
recall = print('Recall:', metrics.recall_score(y_test, redpredictions))
```

```
Accuracy: 0.89475
Precision: 0.8078223080637373
Recall: 0.7902692489371752
```

```
In [53]: pred_prob = redmodel.predict(X_test)
fpr, tpr, _ = metrics.roc_curve(y_test, pred_prob)
auc = metrics.roc_auc_score(y_test, pred_prob)
plt.plot(fpr, tpr, label = 'auc = '+str(auc))
plt.legend(loc = 4)
plt.show()
```



```
In [55]: redmodel.coef_
```

```
Out[55]: array([[ -2.71883489e-01,  7.76387217e-03,  2.24675648e-03,  
                2.37715719e+00,  1.66342876e+00,  1.45445120e+00]])
```

```
In [ ]:
```